CSCI544: Homework Assignment №1

Due on Sept 12, 2023 (before class)

This assignment gives you hands-on experience with text representations and the use of text classification for sentiment analysis. Sentiment analysis is extensively used to study customer behaviors using reviews and survey responses, online and social media, and healthcare materials for marketing and costumer service applications. The assignment is accompanied with a Jupyter Notebook to structure your code. Please submit:

1. A PDF report which contains answers to the questions in the assignment along with brief explanations about your solution. Please also print the completed Jupyter Notebook in PDF format and merge it with your report. If you are comfortable, you can answer to the questions on Jupyter Notebook as well but just submit one PDF file by merging your written answer and the completed Jupyter notebook in case you use a separate PDF to answer the question. On your completed Jupyter notebook, please print the requested values, too.

2. You also need to submit an executable .py file which when run, generates the requested numerical outputs in the assignment as listed at the end of the assignment description. We need the .py file to check overlap between codes to detect plagiarism. Please include the Python version you use.

The libraries that you will need are included in the HW1.ipynb file. You can use other libraries as far as they decently similar to the ones included in the HW1.ipynb file but do not use more advance libraries. If you use online resource, you need to cite and explain how you have used the resource. At the beginning of the .py file, add a read command to read the data.tsv file as the input to your .py file from the current directory.

1. Dataset Preparation (10 points)

We will use the Amazon reviews dataset which contains real reviews for office products sold on Amazon. The dataset is downloadable at:

    https://web.archive.org/web/20201127142707if_/https://s3.amazonaws.
    com/amazon-reviews-pds/tsv/amazon_reviews_us_Office_Products_v1_
    00.tsv.gz

Be patient as it may take some time before you have the dataset download but it will be done is a few minutes.

(a)

Read the data as a Pandas frame using Pandas package and only keep the Reviews and Ratings fields in the input data frame to generate data. Our goal is to train sentiment analysis classifiers.

We create a binary classification problem according to the ratings. Let ratings with the values of 1, 2 and 3 form class 1, and ratings with the values of 4 and 5 form class 2. The original dataset is large. To avoid the computational burden, select 50,000 random reviews from each rating class and create a balanced dataset to perform the required tasks on the downsized dataset. Split your dataset into 80% training dataset and 20% testing dataset. Note that you can split your dataset after step 4 when the TF-IDF features are extracted.

Follow the given order of data processing but you can change the order if it improves your final results.

2. Data Cleaning (20 points)

Use some data cleaning steps to preprocess the dataset you created. For example, you can use:
- convert all reviews into lowercase.
- remove the HTML and URLs from the reviews
- remove non-alphabetical characters
- remove extra spaces
- perform contractions on the reviews, e.g., won't → will not. Include as many contractions in English that you can think of.

You can use other cleaning procedures that can help to improve performance. You can either use Pandas package functions or any other built-in functions. Do not try to implement the above processes manually.

In your report, print the average length of the reviews in terms of character length in your dataset before and after cleaning (to be printed by .py file).

3. Preprocessing (20 points)

Use NLTK package to process your dataset:
    - remove the stop words
    - perform lemmatization
In your report and the .py file, print the average length of the reviews in terms of character length in before and after preprocessing.

4. Feature Extraction (10 points)

Use sklearn to extract both TF-IDF and Bag of Words (BoW) features. Note that BoW may need a little more programming but is not difficult to generate. At this point, you should have created two datasets that consists of features and labels for the reviews you selected.

5. Perceptron (10 points)

Train a Perceptron model on your training dataset using the sklearn built-in implementation.

Report Precision, Recall, and f1-score for training Perceptron using both BoW and TF-IDF features. These 6 values should be printed in two separate lines by the .py file for first BoW and then TF-IDF as follows
    - Precision Recall F1
    - Precision Recall F1

6. SVM (10 points)

Train an SVM model on your training datasets using the sklearn built-in implementation. Report Precision, Recall, and f1-score similar to the previous question format in lines 3 and 4.

7. Logistic Regression (10 points)

Train a Logistic Regression model on your training datasets using the sklearn built-in implementation. Report Precision, Recall, and f1-score similar to the previous question format in lines 5 and 6 by the .py file.

8. Naive Bayes (10 points)

Train a Naive Bayes model on your training dataset using the sklearn built-in implementation. Report Precision, Recall, and f1-score similar to the previous question format in lines 7 and 68 by the .py file.

Note 1: To be consistent, when the .py file is run, the following should be printed, each in a line:

- Average length of reviews before and after data cleaning (with a comma between them)

- Average length of reviews before and after data preprocessing (with comma between them)

- Precision, Recall, and f1-score for the testing split in 2 lines

- Precision, Recall, and f1-score for the testing split in 2 lines

- Precision, Recall, and f1-score for the testing split in 2 lines

- Precision, Recall, and f1-score for the testing split in 2 lines

Note that in your Jupyter notebook, print the Precision, Recall, and f1-score for the above models just by putting a space between them and in .py file in separate lines.

Note 2: Your models should have a decent performance to receive full credit. The decent performance will be determined later by checking all submissions